

Python 3.5.6 |Anaconda 4.2.0 (64-bit)| (default, Aug 26 2018, 16:05:27) [MSC v.1900 64 bit (AMD64)]
Type "copyright", "credits" or "license" for more information.

IPython 5.1.0 -- An enhanced Interactive Python.
? -> Introduction and overview of IPython's features.
%quickref -> Quick reference.
help -> Python's own help system.
object? -> Details about 'object', use 'object??' for extra details.

In [1]:

```
In [1]: import pandas as pd
...: import numpy as np
...: import matplotlib.pyplot as plt
```

```
In [2]: houseData = pd.read_csv('housing.csv')
...:
...: print(houseData.iloc[0:6,:].values)
...: X = houseData.iloc[:, :-1].values
...: y = houseData.iloc[:, 9].values
[[-122.23  37.88  41  880 129.0  322 126  8.3252 'NEAR BAY' 452600]
 [-122.22  37.86  21 7099 1106.0 2401 1138  8.3014 'NEAR BAY' 358500]
 [-122.24  37.85  52 1467 190.0  496 177  7.2574 'NEAR BAY' 352100]
 [-122.25  37.85  52 1274 235.0  558 219  5.6431 'NEAR BAY' 341300]
 [-122.25  37.85  52 1627 280.0  565 259  3.8462 'NEAR BAY' 342200]
 [-122.25  37.85  52  919 213.0  413 193  4.0368 'NEAR BAY' 269700]]
```

```
In [3]: from sklearn.preprocessing import Imputer
...: missingValues = Imputer(missing_values="NaN", strategy="mean", axis=0)
```

```
In [4]: X[:, 0:8] = missingValues.fit_transform(X[:, 0:8])
...:
```

```
In [5]: from sklearn.preprocessing import LabelEncoder
...: X_labelencoder = LabelEncoder()
...:
...: X[:, 8] = X_labelencoder.fit_transform(X[:, 8])
```

```
In [6]: X_labelencoder.classes_
...:
```

```
Out[6]: array(['<1H OCEAN', 'INLAND', 'ISLAND', 'NEAR BAY', 'NEAR OCEAN'],
dtype=object)
```

```
In [7]: from sklearn.preprocessing import OneHotEncoder
...: X_ohe = OneHotEncoder( categorical_features=[8])
...: X = X_ohe.fit_transform(X).toarray()
...:
```

```
In [8]: from sklearn.cross_validation import train_test_split
...: X_train, X_test, y_train, y_test = train_test_split
(X, y, test_size=0.2, random_state=0)
```

```
In [9]: from sklearn.preprocessing import StandardScaler
...: scale = StandardScaler()
...: X_train = scale.fit_transform(X_train)
...: X_test = scale.fit_transform(X_test)
...:
```

```
In [10]: from sklearn.linear_model import LinearRegression
...: linearRegressor = LinearRegression()
...: linearRegressor.fit(X_train, y_train)
```

```

...:
Out[10]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)

In [11]: LRpredict = linearRegressor.predict(X_test)
...:

In [12]: linearRegressor.score(X_train,y_train)
...:
Out[12]: 0.64696546198529004

In [13]: linearRegressor.score(X_test,y_test)
Out[13]: -1.4353247599973323e+21

In [14]: from sklearn.metrics import mean_squared_error
...: from math import sqrt
...: LRrms = sqrt(mean_squared_error(y_test,LRpredict))

In [15]: from sklearn.tree import DecisionTreeRegressor
...: DTregressor = DecisionTreeRegressor(random_state=4)
...: DTregressor.fit(X_train,y_train)
...:
Out[15]:
DecisionTreeRegressor(criterion='mse', max_depth=None, max_features=None,
                      max_leaf_nodes=None, min_samples_leaf=1, min_samples_split=2,
                      min_weight_fraction_leaf=0.0, presort=False, random_state=4,
                      splitter='best')

In [16]: DTpredict = DTregressor.predict(X_test)
...:

In [17]: DTregressor.score(X_train,y_train)
Out[17]: 1.0

In [18]: DTregressor.score(X_test,y_test)
Out[18]: 0.5418211162276273

In [19]: from sklearn.metrics import mean_squared_error
...: from math import sqrt
...: DTrms = sqrt(mean_squared_error(y_test,DTpredict))

In [20]: from sklearn.ensemble import RandomForestRegressor
...: RFregressor = RandomForestRegressor(n_estimators=20,random_state=2)
...: RFregressor.fit(X_train,y_train)

Out[20]:
RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,
                      max_features='auto', max_leaf_nodes=None, min_samples_leaf=1,
                      min_samples_split=2, min_weight_fraction_leaf=0.0,
                      n_estimators=20, n_jobs=1, oob_score=False, random_state=2,
                      verbose=0, warm_start=False)

In [21]: RFPredict = RFregressor.predict(X_test)
...:

In [22]: RFregressor.score(X_train,y_train)
Out[22]: 0.9699115208222413

In [23]: RFregressor.score(X_test,y_test)
Out[23]: 0.7521839910638856

In [24]: from sklearn.metrics import mean_squared_error
...: from math import sqrt

```

```

...: RFrms = sqrt(mean_squared_error(y_test,RFpredict))

In [25]: X = np.delete(X,[0,1,2,3,4,5,6,7,8,9,10,11],axis=1)

In [26]: from sklearn.cross_validation import train_test_split
...: X_train,X_test,y_train,y_test = train_test_split(X,y,
test_size=1/4,random_state=0)
...:

In [27]: X_train.shape
...:
Out[27]: (15480, 1)

In [28]: from sklearn.preprocessing import StandardScaler
...: scaler = StandardScaler()
...: X_train = scaler.fit_transform(X_train)
...: X_test = scaler.fit_transform(X_test)
...:

In [29]: from sklearn.linear_model import LinearRegression
...: linearRegressor1 = LinearRegression()
...: linearRegressor1.fit(X_train,y_train)
Out[29]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)


In [30]: linearRegressor1.predict(X_test)
...:
Out[30]:
array([ 219691.90105431,  287764.77990757,  227926.05134686, ...,
        197854.04791684,  276628.39482425,  245289.27735735])

In [31]: linearRegressor1.score(X_train,y_train)
Out[31]: 0.48061930819884535


In [32]: linearRegressor1.score(X_test,y_test)
...:
Out[32]: 0.45147717106069024

In [33]: plt.scatter(X_train,y_train,color='green')
...: plt.plot(X_train,linearRegressor1.predict(X_train),color='red')
...: plt.title('Median House Price Prediction')
...: plt.xlabel('Median income')
...: plt.ylabel('House price')
...: plt.show()

```



```
In [34]: plt.scatter(X_test,y_test,color='blue')
...: plt.plot(X_train,linearRegressor1.predict(X_train),color='red')
...: plt.title('Median House Price Prediction')
...: plt.xlabel('Median income')
...: plt.ylabel('House price')
...: plt.show()
```



```
In [35]:
```